

# Machine Learning: Basic Principles

## How to Specify A Machine Learning Problem?

Salo, September 2018

## Guiding Questions

- How to formulate your business as a ML Problem ?
- How to determine which algorithm to use for your problem ?

# Outline

- ① Cats or Dogs
- ② From Raw Data to Features and Labels
- ③ Hypothesis Space
- ④ Loss Function
- ⑤ Optimization (=Training)
- ⑥ Wrap Up

# A ML Application




- hard disk full of images
- for a **few images** it is known if they show a cat or a dog
- develop a software tool ( “app” ) to label all images

# Design Choices to be Made

- in which format represent the images ? (what features?)
- which algorithms should we use ? (which hypothesis space?)
- how to evaluate our labelling tool? (how to validate?)
- how to tune for best performance? (how to train?)

# Main Components of a ML Problem

A Machine Learning  
Problem



# Outline

- ① Cats or Dogs
- ② From Raw Data to Features and Labels
- ③ Hypothesis Space
- ④ Loss Function
- ⑤ Optimization (=Training)
- ⑥ Wrap Up

## Raw Data

per se, the dataset is just a (huge) **pile of bits**




clever parsing of data might be most difficult part of ML problem!




## From Raw Data to Vectors (Data Points)

- need to parse raw data into more manageable form
- break raw data into **atomic pieces (data points)**




- $i$ th data point encoded by vector  $\mathbf{z}^{(i)} \in \mathbb{R}^d$
- dataset amounts to a bunch of vectors  $\{\mathbf{z}^{(i)}\}_{i=1}^N$

## From Audio to Vectors



what are typical values of  $\Delta T$  and  $T_{\text{song}}$  for rock song ?

## From RGB Images to Vectors



# Labeled Data

- **partition data point** as  $\mathbf{z} := (\mathbf{x}, y)$
- **input “features”**  $\mathbf{x} \in \mathcal{X}$ , **“label” / “output” / “target”**  $y \in \mathcal{Y}$



data point  $\mathbf{z} = (\mathbf{x}, \text{“dog”})$

image pixels  $\mathbf{x} \in \mathcal{X} = \mathbb{R}^d$

label (ouput)  $y \in \mathcal{Y} = \{\text{“dog”}, \text{“cat”}\}$


- applications with **discrete**  $\mathcal{Y}$  called **classification** problems
- applications with **continuous**  $\mathcal{Y}$  called **regression** problems

## A Key Message

- in real-life applications its not obvious what part of data is label and what should be the features
- feature and label space  $\mathcal{X}, \mathcal{Y}$  are **design choice!** (we have to find the most useful choices for our application at hand!)
- HOWEVER, there are methods to automatically choose good features (“Feature Learning”)

## Labeled Data for Regression


data point  $\mathbf{z}^{(i)}$  consists of pixels  $\mathbf{x}^{(i)}$  (=features) and temperature  $y^{(i)}$  (=label)



$$\mathbf{z}^{(1)} = (\mathbf{x}^{(1)}, y^{(1)} = 8)$$



$$\mathbf{z}^{(2)} = (\mathbf{x}^{(2)}, y^{(2)} = 8)$$



$$\mathbf{z}^{(3)} = (\mathbf{x}^{(3)}, y^{(3)} = 23)$$

# Label Information is worth Gold!


- accurate **label information is extremely precious**
- ML most powerful with vast amounts of labeled data (=training data)
- HOWEVER, obtaining labels is typically costly
- “labelling” of data often requires human (expert) labour

## Aquiring Labels in Marine Biology





## Acquiring Labels in Particle Physics



## Acquiring Labels in Pharmacology



doing a good job as ML scientist/engineer might save lives !!!

# The Amazon Mechanical Turk

## What can you build with Amazon Mechanical Turk?

[Learn more about common use cases below](#)

### Image/Video Processing

MTurk is well-suited for processing images. While difficult for computers, it is a task that is extremely easy for people to do. In the past, companies have used MTurk to:



Tag objects found in an image to improve your search or advertising targeting



Review a set of images to select the best picture to represent a product



Audit user-uploaded images or videos to moderate content



Classify objects found in satellite imagery

### Data Verification and Clean-up

you can hire human labelling workforce!

# Labels are Costly!

## Amazon Mechanical Turk Pricing

The price you (the Requester) pay for a Human Intelligence Task ("HIT") is comprised of two components: the Worker Reward and the Mechanical Turk Fee. The fee you pay Mechanical Turk is based on the amount you pay Workers. Additional fees may apply for using the Masters Qualification and Premium Qualifications.


<b>Worker Reward</b>	You decide how much to pay Workers for each assignment.
<b>Mechanical Turk Fee</b>	20% fee on the reward and bonus amount (if any). HITs with 10 or more assignments will incur an additional 20% fee on the reward you pay Workers. Minimum fee is \$0.01 per assignment or bonus payment.
<b>Additional Fee for using the Masters Qualification</b> ( <a href="#">What are Masters?</a> )	5% of the reward you pay Workers.
<b>Additional Fee per assignment for using Premium Qualifications</b> ( <a href="#">How do I use Premium Qualifications?</a> )	Blogger: \$0.25 Born 1918 to 1960 (Age 55 or older): \$0.50 Born 1961 to 1971 (Age 45-55): \$0.50 Born 1972 to 1981 (Age 35-45): \$0.50 Born 1982 to 1986 (Age 30-35): \$0.50 Born 1987 to 1991 (Age 25-30): \$0.50 Born 1992 to 1999 (Age 18-25): \$0.50 Borrower - Auto Loans: \$0.40

# Outline

- ① Cats or Dogs
- ② From Raw Data to Features and Labels
- ③ Hypothesis Space**
- ④ Loss Function
- ⑤ Optimization (=Training)
- ⑥ Wrap Up

# Hypothesis Maps Features to Labels

- want to predict label  $y \in \mathcal{Y}$  from features  $\mathbf{x} \in \mathcal{X}$  of data point
- consider hypothesis map  $h(\cdot) : \mathcal{X} \rightarrow \mathcal{Y}$
- hypothesis for **discrete**  $\mathcal{Y}$  (e.g.,  $\mathcal{Y} = \{0, 1\}$ ) called **classifier**
- hypothesis for **continuous**  $\mathcal{Y}$  (e.g.,  $\mathcal{Y} = \mathbb{R}$ ) called **predictor**



## How Good is A Predictor?

- ML is about finding good predictor  $h(\cdot) : \mathcal{X} \rightarrow \mathcal{Y}$
- we predict label  $y$  from features  $\mathbf{x}$  by  $\hat{y} = h(\mathbf{x})$
- choose predictor  $h(\cdot)$  such that  $h(\mathbf{x}) \approx y$
- two issues here:
  - i1: set of maps  $h(\cdot) : \mathcal{X} \rightarrow \mathcal{Y}$  is typically LARGE (infinite)
  - i2: need a measure for quality of particular  $h(\cdot)$

# The Hypothesis Space

- GOAL of ML: find predictor  $h(\cdot)$  such that  $h(\mathbf{x}) \approx y$
- two issues here:
  - i1: set of maps  $h(\cdot) : \mathcal{X} \rightarrow \mathcal{Y}$  is typically LARGE
  - i2: how to measure approximation quality  $h(\mathbf{x}) \approx y$
- solve i1 by restricting  $h(\cdot)$  to **subset  $\mathcal{H}$**  of maps  $\mathcal{X} \rightarrow \mathcal{Y}$
- subset  $\mathcal{H}$  referred to as **hypothesis space**




# The Hypothesis Space Picture

# Representing a Hypothesis/Predictor/Classifier

- ML revolves around finding a good predictor  $h(\cdot) \in \mathcal{H}$
- need efficient (computer-friendly) **representation of  $\mathcal{H}$**
- e.g., **binary classification  $\mathcal{Y} = \{0, 1\}$**  with  $|\mathcal{X}| = K$
- what would  $K$  be for  $512 \times 512$  black/white bitmap?
- how many numbers specify an arbitrary map  $\mathcal{X} \rightarrow \mathcal{Y}$ ?

## Representing a Hypothesis via Decision Boundary

- binary classification with  $\mathcal{Y} = \{0, 1\}$  and  $\mathcal{X} = \mathbb{R}^2$



- map  $h(\cdot): \mathcal{X} \rightarrow \{0, 1\}$  characterized by decision boundary (DB)
- hypothesis space defined by **allowed shapes of DB**
- e.g.,  $\mathcal{H} = \{ \text{classifiers with DB consisting of 4 line segments} \}$

# A Regression Problem

- $i$ th snapshot represented by feature vector  $\mathbf{x}^{(i)} \in \mathcal{X} = \mathbb{R}^d$



- what is  $d$  for snapshots being  $512 \times 512$  RGB bitmap?
- we label  $i$ th snapshot by local temperature  $y^{(i)} \in \mathcal{Y} = \mathbb{R}$

## Representing a Hypothesis for Regression ( $\mathcal{Y} = \mathbb{R}$ )

- 512 × 512 RGB webcam snapshot  $\mathbf{x}^{(i)} \in \mathcal{X} \subseteq \mathbb{R}^d$
- snapshot labeled with temperature  $y^{(i)} \in \mathcal{Y} = \mathbb{R}$

- hypothesis space  $\mathcal{H}$  of linear regression:

$$\mathcal{H} = \{h^{(\mathbf{w})}(\mathbf{x}) = \mathbf{w}^T \mathbf{x}, \text{ with some } \mathbf{w} \in \mathbb{R}^d\}$$

- is  $\mathcal{H}$  a proper subset of the set of all maps  $\mathbb{R}^d \rightarrow \mathbb{R}$ ?
- choose  $\mathbf{w}$  such that  $y \approx h^{(\mathbf{w})}(\mathbf{x})$

# Parametrizing the Hypothesis Space of Linear Regression



## Representing a Hypothesis via Code

- $h^{(w)}(\mathbf{x}) := \mathbf{w}^T \mathbf{x} \in \mathbb{R}$  predicts temperature for snapshot  $\mathbf{x}$

```
1 function temp=WhatIsTheTemperature (image,weight)
2
3     temp = weight'*image ;
4 end
```

- think of hypothesis as a (Python/Matlab/...) subroutine
- hypothesis space could be, e.g.,  
$$\mathcal{H} = \{ \text{all python routines with runtime less than 10 sec.} \\ \text{and having as input an image and a tuning parameter} \\ \text{and output a temperature} \}$$

# Representing a Hypothesis via Decision Trees



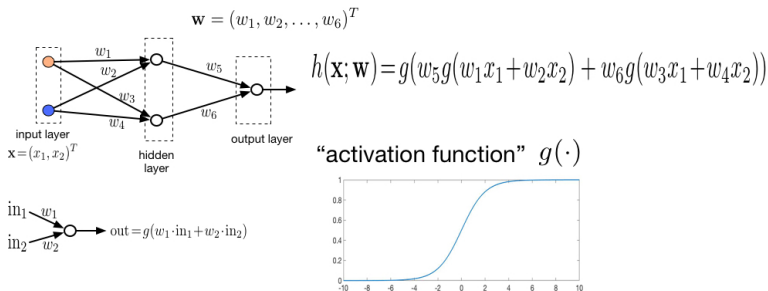
$$\mathcal{R}_1 = \{\mathbf{x} : x_g \geq w, x_r < w\} \quad \mathcal{R}_3 = \{\mathbf{x} : x_g < w\}$$

$$\mathcal{R}_2 = \{\mathbf{x} : x_g \geq w, x_r \geq w\}$$

- fast evaluation of  $h(\mathbf{x})$  by walking down the tree
- e.g.,  $\mathcal{H} = \{ \text{decision trees of depth less than six} \}$



# Representing a Hypothesis via a “Neural Network”



- network representation enables efficient computations !!!
- e.g.,  
 $\mathcal{H} = \{ \text{NN with three hidden layers each having 10 units} \}$

## Representing Hypothesis via Feature Maps

- consider original input vector  $\mathbf{x} \in \mathcal{X} \subseteq \mathbb{R}^d$
- define **feature map**  $\phi(\cdot) : \mathcal{X} \rightarrow \mathcal{F} \subseteq \mathbb{R}^n$  with  $d \ll n$
- high-dimensional feature space  $\mathcal{F}$
- construct non-linear classifiers  $h^{(\mathbf{w})}(\mathbf{x}) := \mathcal{I}(\mathbf{w}^T \phi(\mathbf{x}) > 0)$
- e.g.,  $d = 2$  and  $\phi(\mathbf{x}) = (x_1, x_2, x_1^2 + x_2^2, 1)^T \in \mathbb{R}^4$  ( $n = 4$ )
- what is decision boundary of  $h^{(\mathbf{w})}(\mathbf{x})$  for  $\mathbf{w} = (0, 0, 1, -1)^T$ ?
- feature maps used in kernel methods (see course CS-E4830)


# Linear Classifiers

- binary classification  $\mathcal{Y} = \{0, 1\}$  with feature space  $\mathcal{X} = \mathbb{R}^d$
- classifier  $h(\cdot) : \mathcal{X} \rightarrow \mathcal{Y}$  represented by decision boundary
- how many different decision boundaries are there?
- restrict  $h(\cdot)$  to manageable subset  $\mathcal{H}$  (hypothesis space)
- **linear classifiers** are particular hypothesis space

$$\mathcal{H} := \{h(\cdot) \text{ with decision boundary being hyperplane} \}$$

## Linear Binary Classifiers for $\mathcal{X} = \mathbb{R}^2$

decision boundaries are straight lines



how many linear classifiers do exist for  $\mathcal{X} = \mathbb{R}^2$  ?

# Outline

- ① Cats or Dogs
- ② From Raw Data to Features and Labels
- ③ Hypothesis Space
- ④ Loss Function**
- ⑤ Optimization (=Training)
- ⑥ Wrap Up

# The Quality of a Hypothesis

- GOAL of ML: choose hypothesis  $h(\cdot)$  such that  $h(\mathbf{x}) \approx y$
- two issues here:
  - i1: set of maps  $h(\cdot) : \mathcal{X} \rightarrow \mathcal{Y}$  is typically LARGE
  - i2: how to measure approximation quality  $h(\mathbf{x}) \approx y$
- i2 requires measure for loss/error incurred by predictor  $h(\mathbf{x})$
- define loss function  $L(\mathbf{z}, h(\cdot))$  incurred by  $h(\cdot)$  for data point  $\mathbf{z}$
- most reasonable loss functions share structural similarities  
<http://web.mit.edu/lrosasco/www/publications/loss.pdf>


# The Squared-Error Loss

- consider labeled data  $\mathbb{X} = \{\mathbf{x}^{(i)}, y^{(i)}\}_{i=1}^N$
- **continuous labels**  $y^{(i)} \in \mathbb{R}$  (**regression** problem)
- we predict label  $y^{(i)}$  using predictor  $h(\mathbf{x}^{(i)})$
- natural choice is **squared error**  $L((\mathbf{x}, y), h(\cdot)) := (y - h(\mathbf{x}))^2$



## The 0/1 Loss

- consider labeled data  $\mathbb{X} = \{\mathbf{x}^{(i)}, y^{(i)}\}_{i=1}^N$
- **binary labels**  $y^{(i)} \in \{-1, 1\}$  (**classification** problem)
- we predict label  $y^{(i)}$  using predictor  $h(\mathbf{x}^{(i)})$
- natural choice is **0/1-loss**  $L((\mathbf{x}, y), h(\cdot)) := \mathcal{I}(yh(\mathbf{x}) > 0)$






# The Empirical Risk

- consider a particular loss function  $L(\mathbf{z}, h(\cdot))$
- evaluate loss for data points in the dataset  $\mathbb{X}$
- empirical/training loss/risk/error

$$\mathcal{E}(h(\cdot)|\mathbb{X}) := (1/N) \sum_{i=1}^N L((\mathbf{x}^{(i)}, y^{(i)}), h(\cdot))$$

- $\mathcal{E}(h(\cdot)|\mathbb{X})$  is **mean squared error** for squared error loss
- $\mathcal{E}(h(\cdot)|\mathbb{X})$  is **misclassification rate** for 0/1 loss

# Mean Squared Error



$$\mathbf{z}^{(1)} = (\mathbf{x}^{(1)}, y^{(1)} = 8)$$


$$h(\mathbf{x}^{(1)}; \mathbf{w}) = 10$$



$$\mathbf{z}^{(2)} = (\mathbf{x}^{(2)}, y^{(2)} = 8)$$

$$h(\mathbf{x}^{(2)}; \mathbf{w}) = 7$$

$\mathbf{x}^{(3)} \left\{ \right.$



$$\mathbf{z}^{(3)} = (\mathbf{x}^{(3)}, y^{(3)} = 23)$$

$$h(\mathbf{x}^{(3)}; \mathbf{w}) = 20$$

$y^{(3)}$

$$\mathcal{E}(h|\mathbb{X}) = (1/N) \sum_{i=1}^N (h(\mathbf{x}^{(i)}) - y^{(i)})^2 = (1/3)(2^2 + 1^2 + 3^2) = 14/3$$

# Multitask Learning

- might have to solve different tasks using same dataset
- consider, e.g., dataset of webcam snapshots
- Task 1: predict local temperature  $y$  using pixels  $\mathbf{x}$
- Task 2: classify img into winter/summer using pixels  $\mathbf{x}$
- **individual loss function** for each task:  $\mathcal{E}_1(h(\cdot)|\mathbb{X}), \mathcal{E}_2(h(\cdot)|\mathbb{X})$
- choose  $h(\cdot)$  to **balance optimally** between these two

# Outline

- ① Cats or Dogs
- ② From Raw Data to Features and Labels
- ③ Hypothesis Space
- ④ Loss Function
- ⑤ Optimization (=Training)
- ⑥ Wrap Up

## Finding Optimal Hypothesis

- ML amounts to finding **best predictor/classifier**, i.e.,

$$\hat{h} = \underset{h(\cdot) \in \mathcal{H}}{\operatorname{argmin}} \mathcal{E}(h|\mathbb{X})$$

- solution of this **empirical risk minimization** yields two things:
  - “best” classifier/predictor  $\hat{h}(\cdot)$  out of  $\mathcal{H}$
  - minimum empirical error  $\mathcal{E}(\hat{h}|\mathbb{X})$  achievable for  $\mathcal{H}$
- if  $\mathcal{E}(\hat{h}|\mathbb{X})$  is too high then we might enlarge  $\mathcal{H}$
- $\mathcal{E}(\hat{h}|\mathbb{X})$  is indicator for accuracy of predicting  $y$  from  $\mathbf{x}$  using  $\hat{h}(\mathbf{x})$  for a “new” data point (with unknown label  $y$ )

## Optimal Linear Regression

- regression problem  $\mathcal{X} = \mathbb{R}^d$ ,  $\mathcal{Y} = \mathbb{R}$  with squared error loss
- predict label  $y$  from features  $\mathbf{x}$  using predictors

$$\mathcal{H} = \{h^{(\mathbf{w})}(\mathbf{x}) = \mathbf{w}^T \mathbf{x} \text{ with } \mathbf{w} \in \mathbb{R}^d\}$$


- optimal predictor  $\hat{h} = \operatorname{argmin}_{h(\cdot) \in \mathcal{H}} \mathcal{E}(h|\mathbb{X})$
- equivalent to find optimum weight vector  $\mathbf{w}_0$ , i.e.,

$$\mathbf{w}_0 = \operatorname{argmin}_{\mathbf{w} \in \mathbb{R}^d} f(\mathbf{w}) := \mathcal{E}(h^{(\mathbf{w})}|\mathbb{X})$$

- optimum  $\mathbf{w}_0$  characterized by **zero-gradient condition**

$$\nabla f(\mathbf{w}_0) = \mathbf{0}$$

## Zero Gradient is Necessary for Optimum



# Gradient Descent

- best predictor  $h^{(\mathbf{w}_0)}$  obtained for optimal weight

$$\mathbf{w}_0 = \underset{\mathbf{w}}{\operatorname{argmin}} f(\mathbf{w}) \text{ with } f(\mathbf{w}) = \mathcal{E}(h^{(\mathbf{w})}|\mathbb{X})$$

- for convex  $f(\mathbf{w})$ , minimum  $\mathbf{w}_0$  characterized by  $\nabla f(\mathbf{w}) = \mathbf{0}$
- this is equivalent to fixed-point equation  $\mathcal{T}\mathbf{w} = \mathbf{w}$
- here,  $\mathcal{T}\mathbf{w} = \mathbf{w} - \alpha \nabla f(\mathbf{w})$  with step-size (learning rate)  $\alpha > 0$
- find fixed-point  $\mathbf{w}_0$  by fixed-point iterations

$$\mathbf{w}^{(k+1)} = \mathcal{T}\mathbf{w}^{(k)} = \mathbf{w}^{(k)} - \alpha \nabla f(\mathbf{w}^{(k)})$$

- known as **gradient descent** (GD)



# Gradient Descent Picture



# Outline

- ① Cats or Dogs
- ② From Raw Data to Features and Labels
- ③ Hypothesis Space
- ④ Loss Function
- ⑤ Optimization (=Training)
- ⑥ Wrap Up

## To Sum Up

- break raw data into **atomic “data points”**
- representing data points using **features and labels**
- various representations of **predictor/classifier** (decision boundaries, decision trees, neural networks, ...)
- concept of **loss functions and empirical risk**
- choosing optimal predictor by **empirical risk minimization**

# The Nitty-Gritty Details

- how to efficiently solve ERM ?
- how to validate a predictor/classifier ?
- how to choose hypothesis space  $\mathcal{H}$  ?
- what to do if we do not have any labels?

## Design Choices

- loss function and hypothesis space are design choices
- squared error loss + linear hyp. space results in convex opt. problem
- choices guided by computational infrastructure
- if all you have is a spreadsheet app, then deep neural nets might not be the right choice for hypothesis space